

RAPIDS

21.12 Release



[@RAPIDSai](https://twitter.com/RAPIDSai)



<https://github.com/rapidsai>



<https://rapids-goai.slack.com/join>

RAPIDS

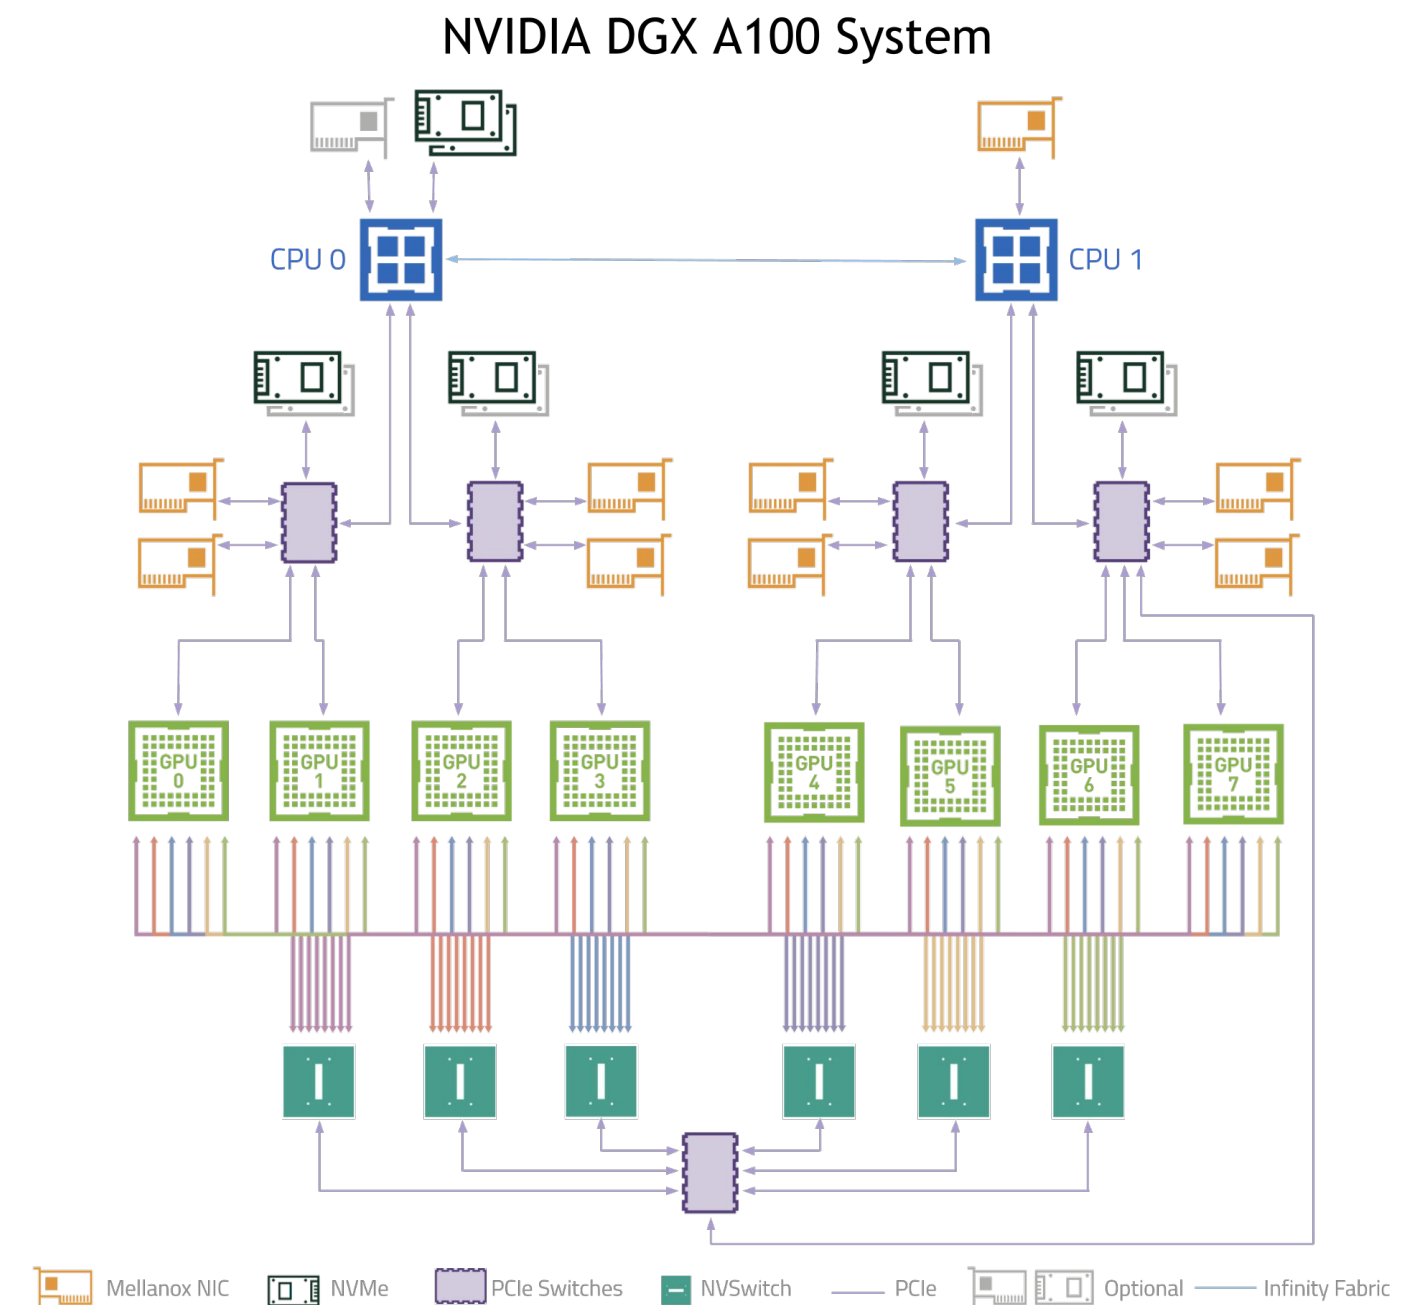
<https://rapids.ai>

Why GPUs for Data Science?

Numerous hardware advantages

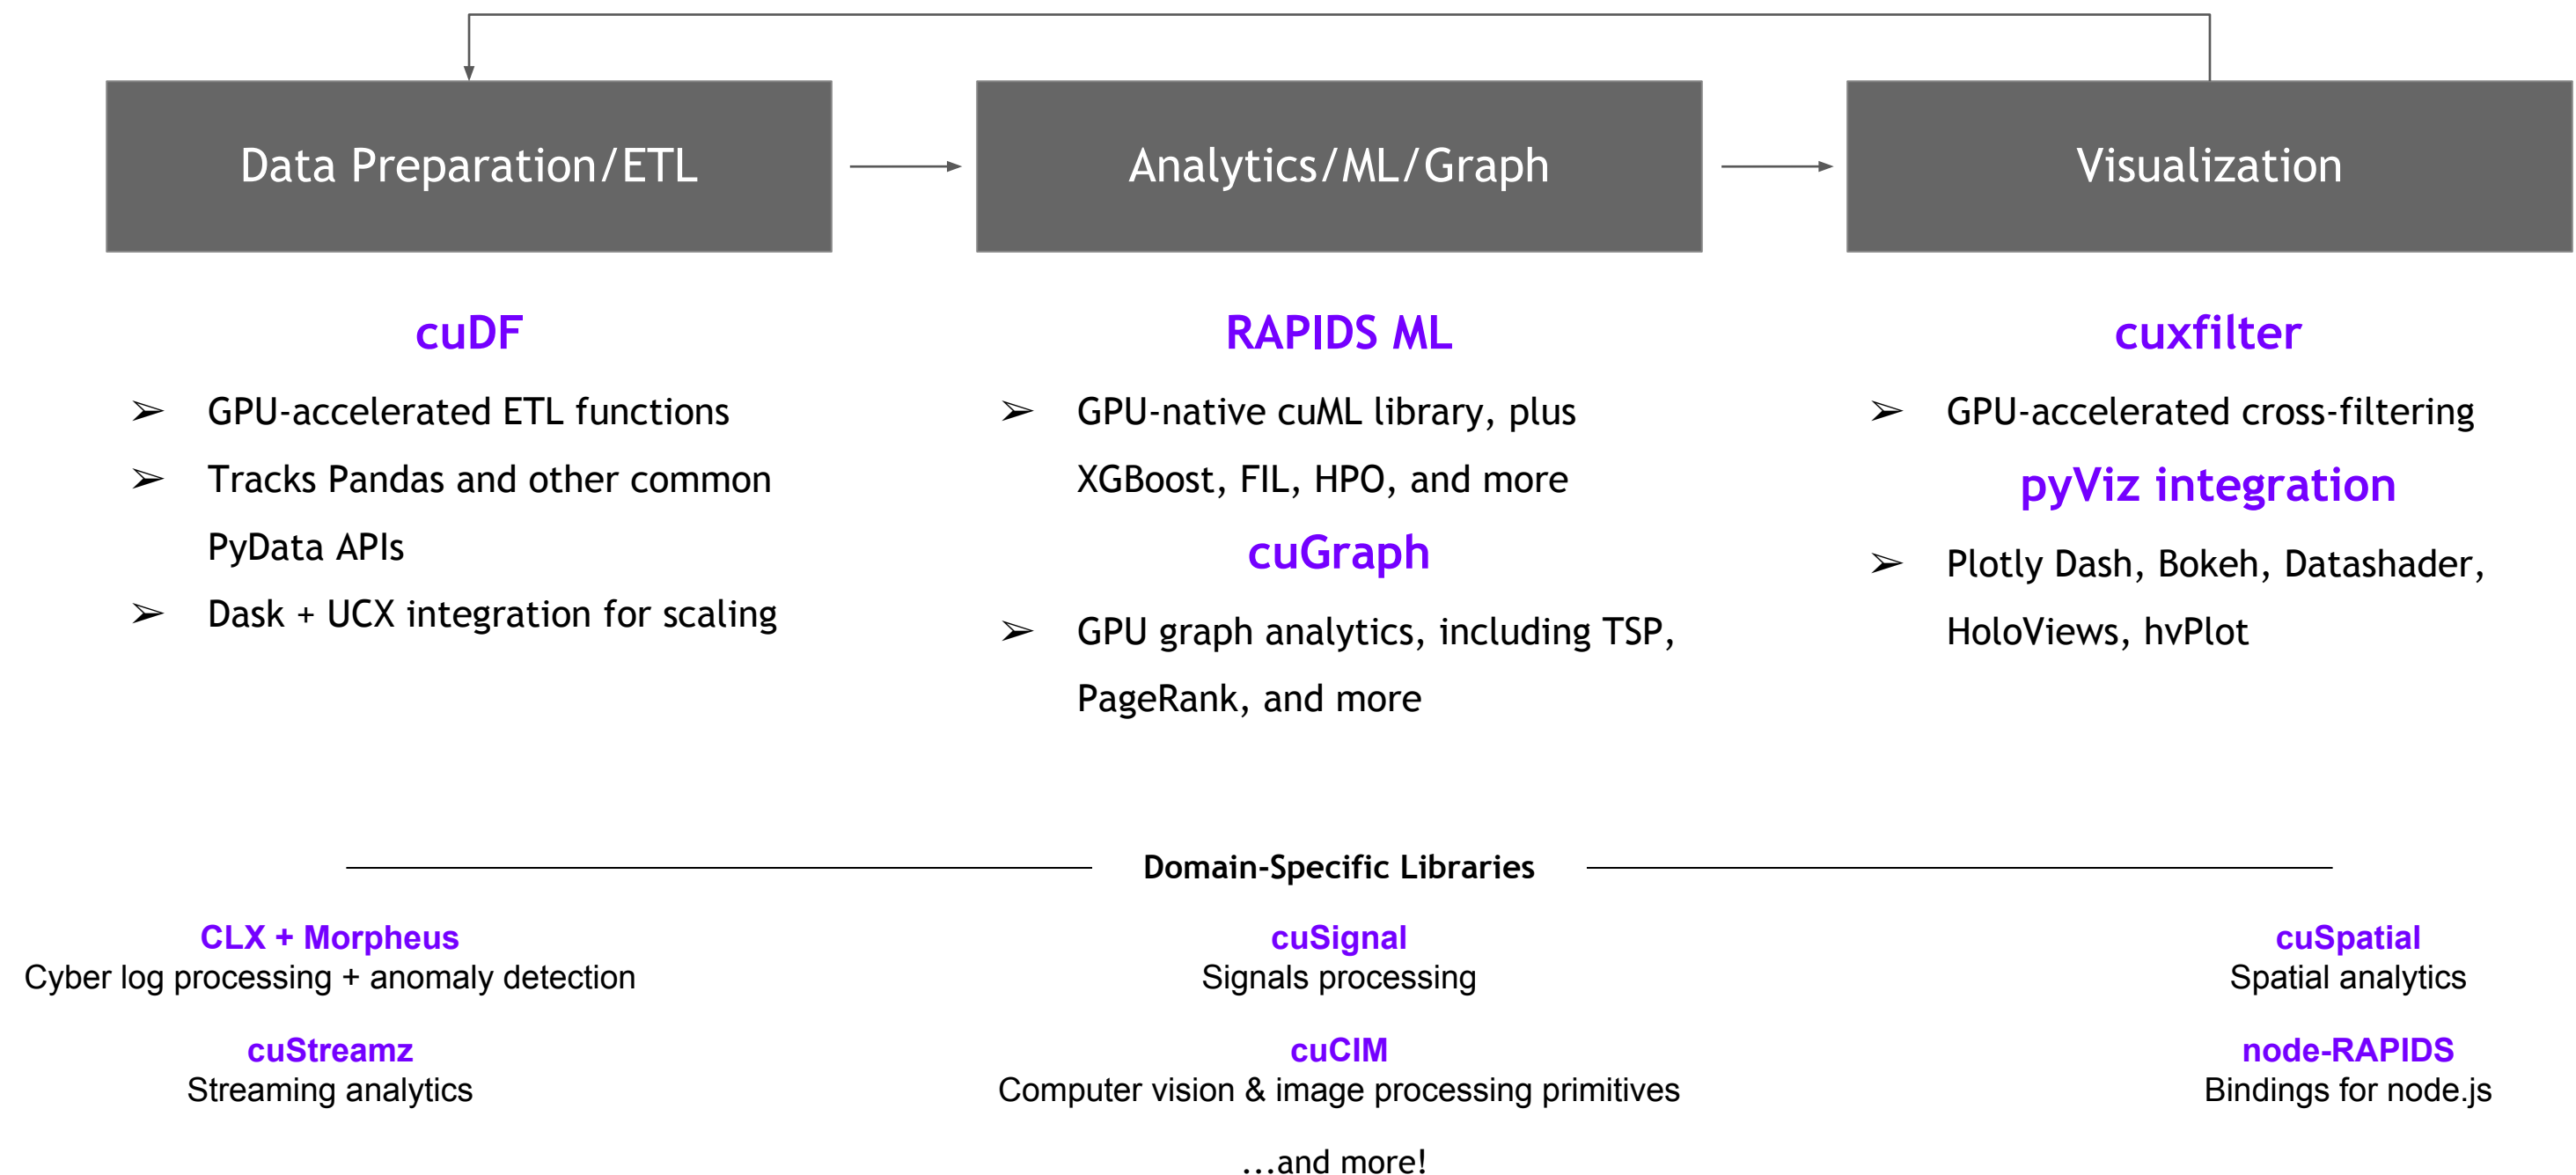
- ▶ Thousands of cores with up to ~20 TeraFlops of general purpose compute performance
- ▶ Up to 1.5 TB/s of memory bandwidth
- ▶ Hardware interconnects for up to 600 GB/s bidirectional GPU <--> GPU bandwidth
- ▶ Can scale up to 16x GPUs in a single node

Almost never run out of compute relative to memory bandwidth!



What is RAPIDS?

End-to-End GPU Accelerated Data Science



Overview of Changes: RAPIDS 21.12 Release

- ▶ **RAPIDS** CUDA enhanced compatibility is now supported by RAPIDS-users can now use newer CUDA binaries or runtimes without updating their CUDA driver version (450.80.02+).
- ▶ **RAPIDS+Dask** Support for spilling on demand; Deprecation of support for UCX version less than 1.11.1
- ▶ **cuDF** 128 bit Decimal support; Groupby diff; Map support for ORC writer; String now has support for istitle; mixed numeric type support in to_dlpack; Resample function added to cuDF; pandas-like UDF interface with DataFrame and Series.apply;
- ▶ **cuML** New LinearSVC and LinearSVR models; Dask-SQL support for cuML models; exogenous variable support in ARIMA; KL divergence exposed in TSNE;
- ▶ **cuGraph** K-core for undirected graphs for single and multiple GPUs; graph primitive filter for edges; transpose_edgelist and symmetrize_edgelist and count number of self loops and multi-edges in a graph in libcugraph; Improved performance for NetworkX inputs; Initial implementation of Multi GPU HITS;
- ▶ **Dask-SQL** Experimental GPU support for Dask-SQL, bringing the power of RAPIDS to SQL

cuDF Updates: Deep Dive

Release 21.12

Features added in 21.12

- ▶ 128 bit Decimal support in libcudf
- ▶ [Diff](#) is now supported for groupby
- ▶ Map support for ORC writer
- ▶ Support for Series.str.[istitle](#)
- ▶ pandas-like UDF interface with DataFrame and Series.[apply](#)
- ▶ Handling of mixed numeric types supported in [to_dlpack](#)
- ▶ [Resample](#) function added to cuDF will allow users to resample time series data
- ▶ cuDF's [Grouper](#) function now supports grouping via time frequency

Planned Upcoming Features

- ▶ Expanded support for 128 bit decimal types
- ▶ Nested type support for JSON reader
- ▶ Decimal support for Dask cuDF Parquet reader
- ▶ Refactored hash join implementation

cuML Updates: Deep Dive

Release 21.12

Features added in 21.12

- ▶ Support for LinearSVM using QN solvers, [LinearSVC](#) and [LinearSVR](#) models
- ▶ Ability to use Dask-SQL with cuML models
- ▶ [ARIMA](#) now supports exogenous variables
- ▶ KL divergence exposed in *TSNE*

Planned Upcoming Features

- ▶ Multiple improvements to T-Distributed Stochastic Neighbor Embedding algorithm to improve accuracy
- ▶ Enhancements to Random Forest Algorithm
- ▶ Symbolic classification and regression models

cuGraph Updates: Deep Dive

Release 21.12


Features added in 21.12

- ▶ Implemented K-core for undirected graphs for Single and Multi GPUs
- ▶ Initial implementation of Multi GPU HITS
- ▶ Initial version of node2vec in C++/CUDA for graph sampling, Python wrapper coming soon
- ▶ New graph filter for the edges, `extract_if_e`, in `libcugraph`
- ▶ Implemented `transpose_edgelist` and `symmetrize_edgelist` in `libcugraph`
- ▶ Added the ability to count number of self loops and multi-edges in a graph, in `libcugraph`
- ▶ Improved performance when a NetworkX graph is passed in

Planned Upcoming Features


- ▶ *Multi-GPU graph primitives for Triangle Counting*
- ▶ *More algorithms ported to using primitives and scaling to Multi GPUs*
- ▶ *Ability for user to write custom algorithms using the graph primitives*
- ▶ *Multi GPU neighborhood sampling*
- ▶ *Property Graphs*

Join the Conversation




GOOGLE GROUPS

<https://groups.google.com/forum/#!forum/rapidsai>



DOCKER HUB

<https://hub.docker.com/r/rapidsai/rapidsai>



SLACK CHANNEL

<https://rapids-goai.slack.com/join>



STACK OVERFLOW

<https://stackoverflow.com/tags/rapids>

THANK YOU



RAPIDS