RAPIDS

22.06 Release

@RAPIDSai

https://github.com/rapidsai

https://rapids-goai.slack.com/join

RAPIDS

https://rapids.ai

Why GPUs for Data Science?

Numerous hardware advantages

- Thousands of cores with up to ~20 TeraFlops of general purpose compute performance
- Up to 1.5 TB/s of memory bandwidth
- Hardware interconnects for up to 600 GB/s bidirectional GPU <--> GPU bandwidth
- Can scale up to 16x GPUs in a single node

Almost never run out of compute relative to memory bandwidth!

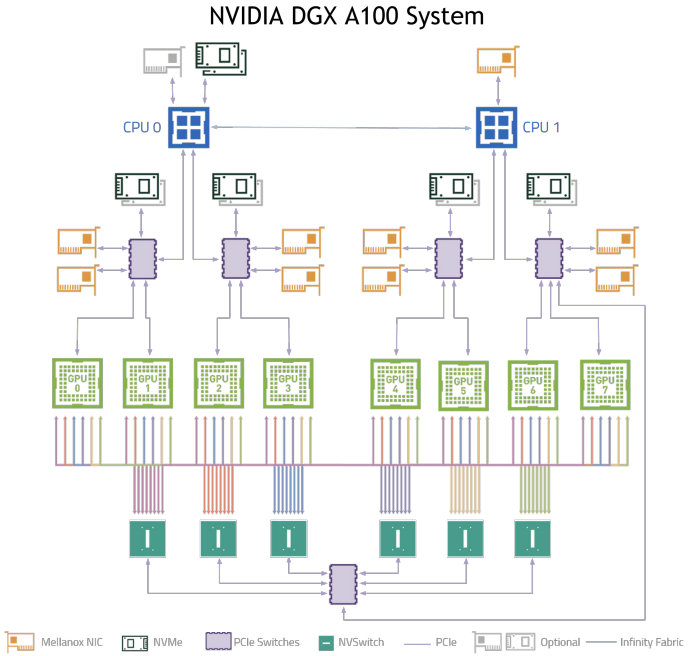

What is RAPIDS?

End-to-End GPU Accelerated Data Science

cuDF

- GPU-accelerated ETL functions
- Tracks Pandas and other common PyData APIs
- Dask + UCX integration for scaling

RAPIDS ML

- GPU-native cuML library, plus XGBoost, FIL, HPO, and more

cuGraph

- GPU graph analytics, including TSP, PageRank, and more

cuxfilter

- GPU-accelerated cross-filtering

pyViz integration

- Plotly Dash, Bokeh, Datashader, HoloViews, hvPlot

Domain-Specific Libraries

CLX + Morpheus
Cyber log processing + anomaly detection

cuStreamz
Streaming analytics

cuSignal
Signals processing

cuCIM
Computer vision & image processing primitives

...and more!

cuSpatial
Spatial analytics

node-RAPIDS
Bindings for node.js

Major updates from the RAPIDS 22.06 release

- **The release of a new RAPIDS Graph-as-a-Service library** GaaS uses cuGraph, cuDF, and other libraries on the server to execute graph data prep and analysis on server-side GPUs. Multiple clients can connect to the server allowing different users and processes the ability to access large graph data that may not otherwise be possible

- **The addition of Multi-Node Multi-GPU (MNMG) algorithms across many libraries, including cuML and cuGraph** When working with massive amounts of data, it's not uncommon to run out of room on a single GPU. We continue to expand support for Multi-Node Multi-GPU (MNMG) workloads across RAPIDS libraries

- **Support for interactive debugging and profiling of memory usage with RMM** Memory allocations in RAPIDS are performed by the RAPIDS Memory Manager (RMM). In the 22.06 release of RMM, users can now trigger a callback function when RAPIDS libraries allocate or free memory for interactive debugging, profiling, and a wide range of other use cases

Overview of Changes: RAPIDS 22.06 Release

- **cuDF** Addition of df.eval() for faster and more memory efficient calculations, df.applymap() for element-wise functions
- **cuML** Initial support for multi-node Logistic Regression, based on Dask-GLM, and Dask's CuPy array support
- **cuGraph** Property Graph available on MNMG architectures; eigenvector centrality algorithm, and Triangle counting algorithm on data located across multiple GPUs
- **Dask-CUDA** Users can now trigger a callback function when RAPIDS libraries allocate or free memory for interactive debugging, profiling, and a wide range of other use cases

cuDF Updates: Deep Dive

Release 22.06

Features added in 22.06

- Decompression by “Zstandard” released for Parquet and ORC (experimental)
- Merging/Joining is 10-30% faster, due to performance optimization
- Documentation updates to user guide, new contributors, library design, C++ doxygen
- Upgrade cudf to support pandas 1.4.x versions
- Add a section to the docs that compares cuDF with Pandas
- Add cudf.DataFrame.applymap
- Implement DataFrame.eval using libcudf ASTs

Planned Upcoming Features

- Support

cuML Updates: Deep Dive

Release 22.06

Features added in 22.06

- First version of Multi-Node Multi-GPU Logistic Regression
- LabelEncoder functionality for CuPy and NumPy arrays
- Random Ball Cover support for 3D data in Nearest Neighbors
- Some other minor tidbits: sample_weight support for Ridge Regression, simplicial_set functions in UMAP
- Double precision support in FIL

Planned upcoming features

cuGraph Updates: Deep Dive

Release 22.06

Features added in 22.06

- Released the initial version of Graph-as-a-Service
- Initial version of a Multi-GPU Property Graph
- Initial release of Python wrappers around cugraph-ops aggregators (pylibcugraphops)
- New Degree Centrality algorithm
- New Multi-GPU Eigenvector Centrality algorithm
- New Multi-GPU Triangle Counting algorithm
-

Planned Upcoming Features

- Data Masking
- Integration with DGL

Join the Conversation

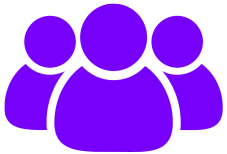

GOOGLE GROUPS

https://groups.google.com/forum/#!forum/rapidsai

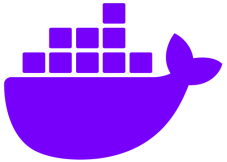

DOCKER HUB

https://hub.docker.com/r/rapidsai/rapidsai

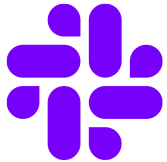

SLACK CHANNEL

https://rapids-goai.slack.com/join

STACK OVERFLOW

https://stackoverflow.com/tags/rapids

THANK YOU

@RAPIDSai

RAPIDS